

PROJECT REPORT



TIME SERIES FORECASTING (SUNSPOTS)

SUBMITTED BY: RATNESH DUBEY (UE178074)

PRAYANSH PANDEY (UE178068)

MD SHAHNAWAZ ALAM (UE178058)

BRANCH - IT - SEM 8TH

SUBMITTED TO: Dr. VEENU MANGAT

TABLE OF CONTENTS

| SNO | Contents |
|------------|---------------------------------------|
| 1 | <u>Introduction</u> |
| 2 | <u>Motivation</u> |
| 3 | <u>Objective</u> |
| 4 | <u>Model Components</u> |
| 5 | <u>Technology Used</u> |
| 6 | <u>Dependencies</u> |
| 7 | <u>Data Sources</u> |
| 8 | <u>Data Overview</u> |
| 9 | <u>Proposed Methodology</u> |
| 10 | <u>Work Plan</u> |
| 11 | <u>Objectives Achieved</u> |
| 12 | <u>Results</u> |
| 13 | <u>Code</u> |
| 14 | <u>Bibliography/References</u> |
| 15 | <u>Future Work</u> |

INTRODUCTION

The sunspot numbers are the major targets which describes the solar activity level. The solar activity influences the human health and living environment on earth. Long-term prediction of sunspot activity can provide important reference information for aerospace, communication, disaster prevention, and so on.

Therefore, the research of sunspot number time-series prediction method has been an important and hot topic in this field. The data of annual and monthly sunspot activity is considered as typical nonlinear, non-Gaussian, and nonstationary time series, which has obvious characteristics of chaotic sequences. The accurate prediction of sunspot numbers has great research significance. It is difficult to predict using the traditional method and the accuracy is not high. [4] [5]

Time series prediction (forecasting) has experienced dramatic improvements in predictive accuracy as a result of the data science machine learning and deep learning evolution. As these ML/DL tools have evolved, businesses and financial institutions are now able to forecast better by applying these new technologies to solve old problems. In this article, we showcase the use of a special type of **Deep Learning model called an LSTM (Long Short-Term Memory)**, which is useful for problems involving [sequences with autocorrelation](#). We analyze a famous historical data set called “[sunspots](#)” (a [sunspot](#) is a solar phenomenon wherein a dark spot forms on the surface of the sun). We’ll show you how you can use an LSTM model to predict sunspots 11 years into the future with an LSTM model.

We have also implemented some other models to compare their performance with LSTM model.

MOTIVATION

The variability of solar activity over time demonstrates highly complex behavior. Neural networks provide useful models for predicting that behavior. Accurate predictions of solar activity are essential for satellite tracking and constellation management, orbit lifetime studies, and mission

planning for long-term orbiting platforms, such as the proposed space station Freedom.

OBJECTIVE

The aim of this project was to develop a Neural Network model for forecasting sunspot activity accurately, we also want to predict the behavior of SC 25. The simplicity and recent success of Deep Neural Networks in forecasting solar-terrestrial phenomena explain why this method has been chosen for the SC 25 prediction.

MODEL COMPONENTS

1. Neural Networks
2. Convolution Layers
3. RNN Layers
4. LSTM Layers
5. Lambda Layers
6. Callbacks
7. Dense Layers.

TECHNOLOGY USED





DEPENDENCIES


- Internet Connection
- GPU(Graphical Processing Unit)
- Labelled Data.
- External Libraries

DATA SOURCES

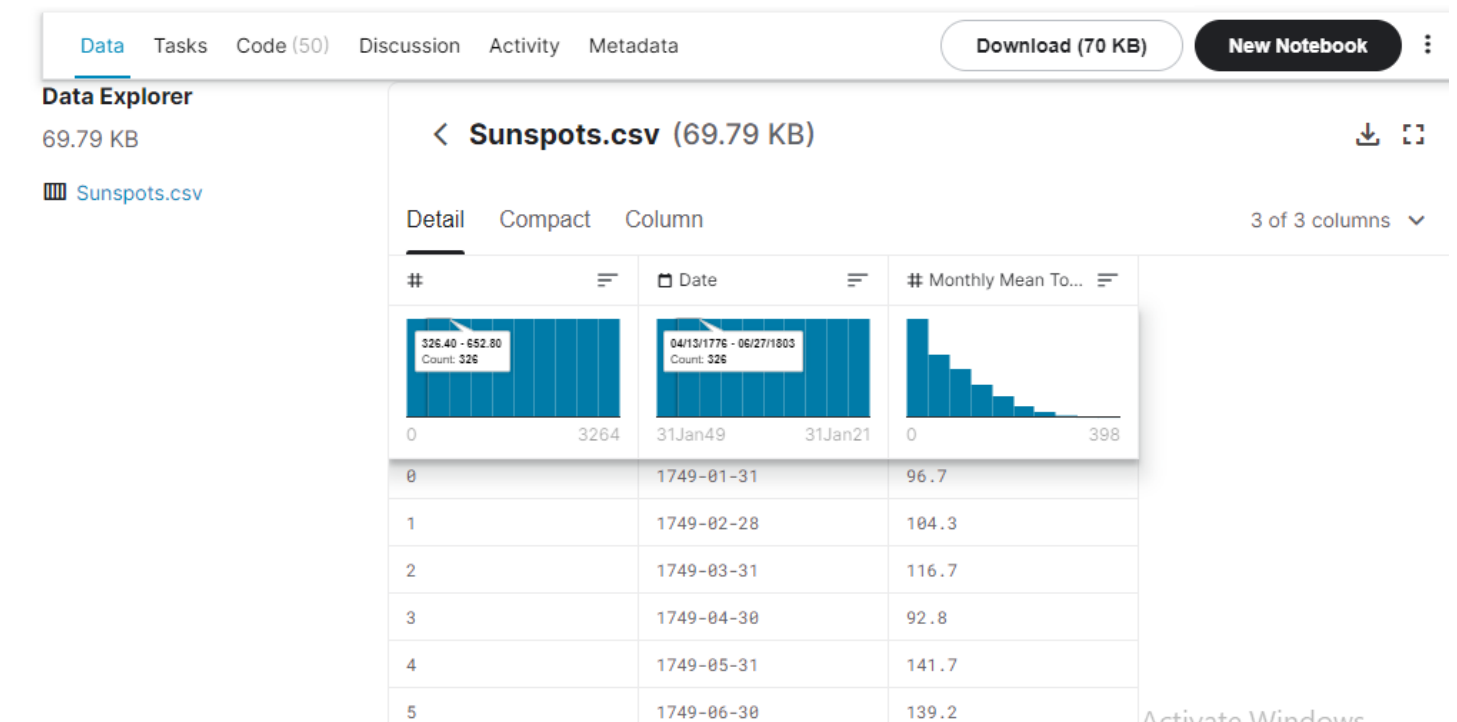
- Kaggle Dataset (Link: <https://www.kaggle.com/robervalt/sunspots>)

DATA OVERVIEW

[Sunspots](#) is a famous data set that ships with R (refer to the Datasets library). The dataset tracks sunspots, which are the occurrence of a dark spot on the sun. Here's an image from NASA showing the solar phenomenon. Pretty cool!

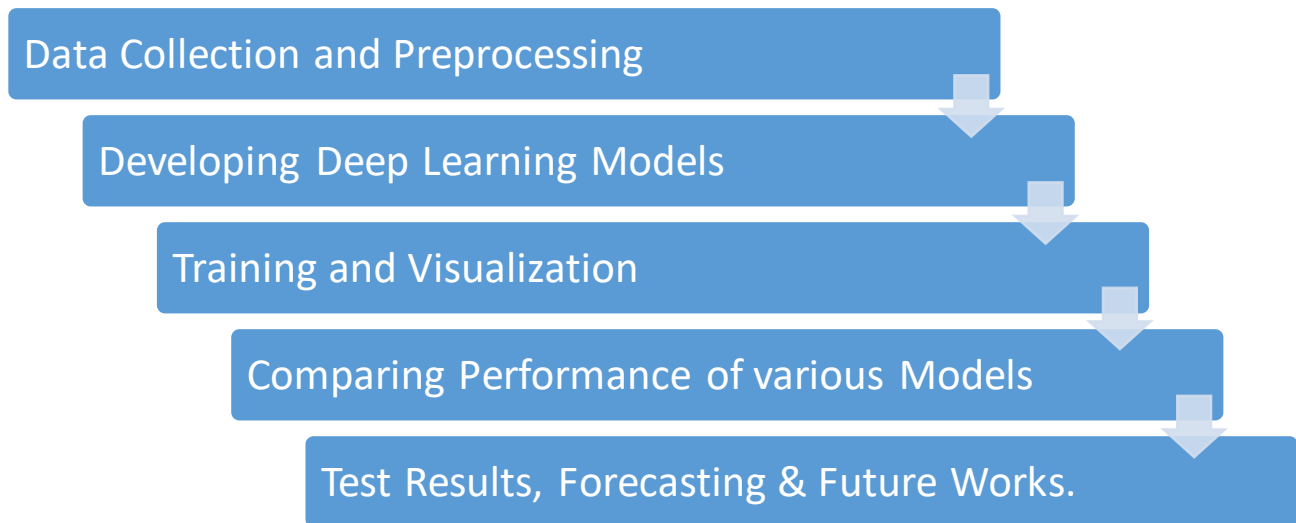


The dataset is a CSV file which contains 3 columns and 3236 rows , 1st column is the serial number, 2nd is date and the third is Monthly Mean Total Sunspot Number.



PROPOSED METHODOLOGY


The methodology will contain the following steps:



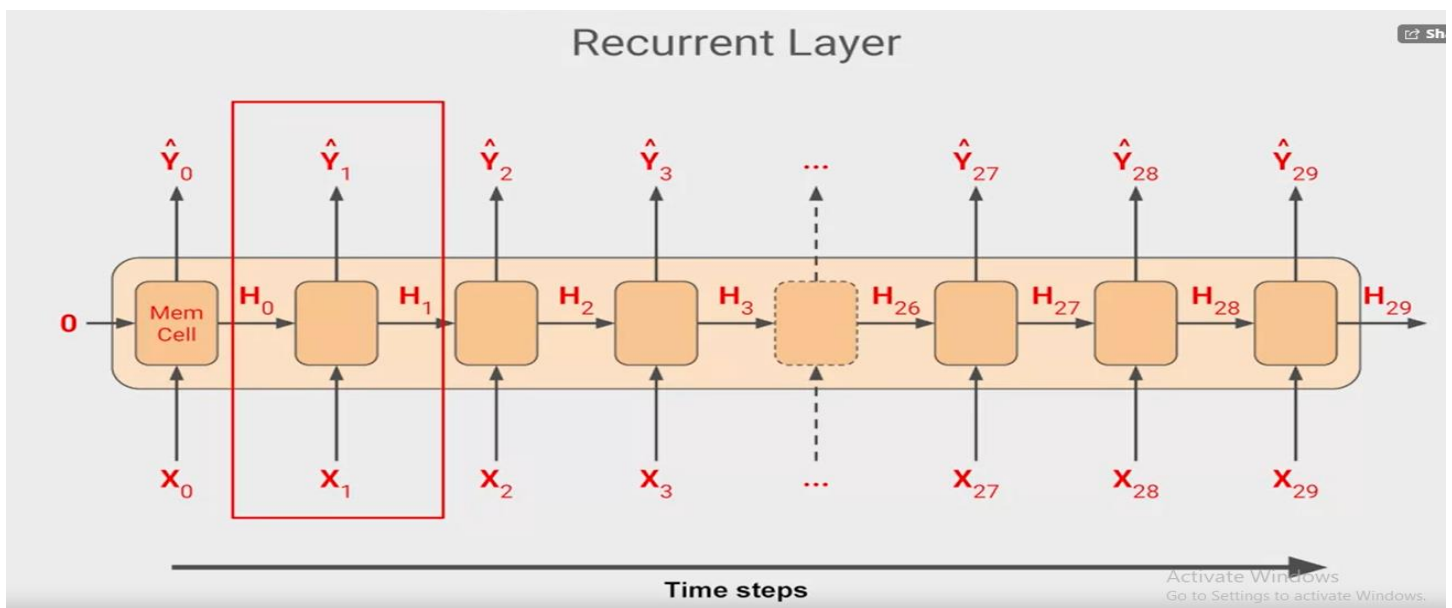
WORK PLAN

- 1.** First we will import the Sunspots data set and perform some preprocessing steps on it.
- 2.** After importing and preprocessing of data, we will do some exploratory data analysis on the data, we try to plot some graphs to see how data trends over time.
- 3.** Create train/test/validation dataset from original data.
- 4.** After that we will apply different models on the data, these contain naïve approach, Moving Averages, Linear Regression, RNNs, and LSTM & Convolutions.
- 5.** We can also add some dense and lambda layers to the model for multilayer Neural Network training, normalization and scaling purpose.
- 6.** Plotting and predicting the results and calculating the mean absolute error.
- 7.** We will then compare the performances of the various models.
- 8.** Finally whole project will be implemented in the form of interface system for users.

A simple LSTM network.



LSTM are just like RNNs [8]



OBJECTIVES ACHIEVED

1. Six Models are created.
2. Performance and loss is compared
3. Graphs are plotted.
4. Prediction is made.
5. LSTM combined with Convolutions seems to perform best with this data.

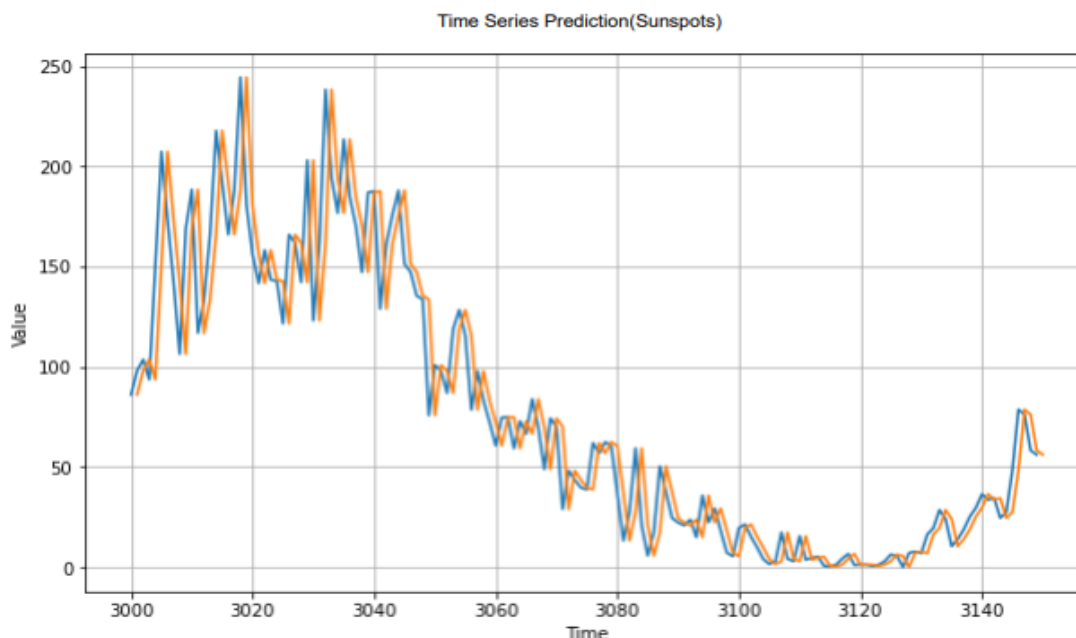
RESULTS

We have used 6 different models to forecast our Time Series, and in each model we compute the mean absolute error, as well as plotted the graph to better understand how well the model is performing.

Naïve Approach [2] [8]

In Naïve Approach we used the most basic approach of using the previous outputs as the future outputs.

$V(t) = V(t-1)$ //here t is current time, and $t-1$ is one time step before current time.



Here we got the Mean Absolute Error of around 16.14, we will take this as a baseline.


Moving Averages & Differencing [2] [8]

In statistics, a moving average is a calculation used to analyze data points by creating a series of averages of different subsets of the full data set.

In layman terms Moving Averages is a statistical technique in which values are averaged over a period of time to predict future values at particular time.

Moving averages can be used to reduce the noise in the data but it does not consider seasonality & trend, so as to prevent this, we first use differencing on the series to remove seasonality and trend, this can be done by subtracting the value at some previous time from current time.

After performing this step, we apply moving averages on the series and then again add the subtracted values in the series to get our forecasted series.



Mean Absolute and Mean Sqaure Error for Moving Averages& Differencing

```
print(tf.keras.metrics.mean_squared_error(x_valid, diff_moving_avg_plus_past).numpy())  
print(tf.keras.metrics.mean_absolute_error(x_valid, diff_moving_avg_plus_past).numpy())
```

2976.2167495035455

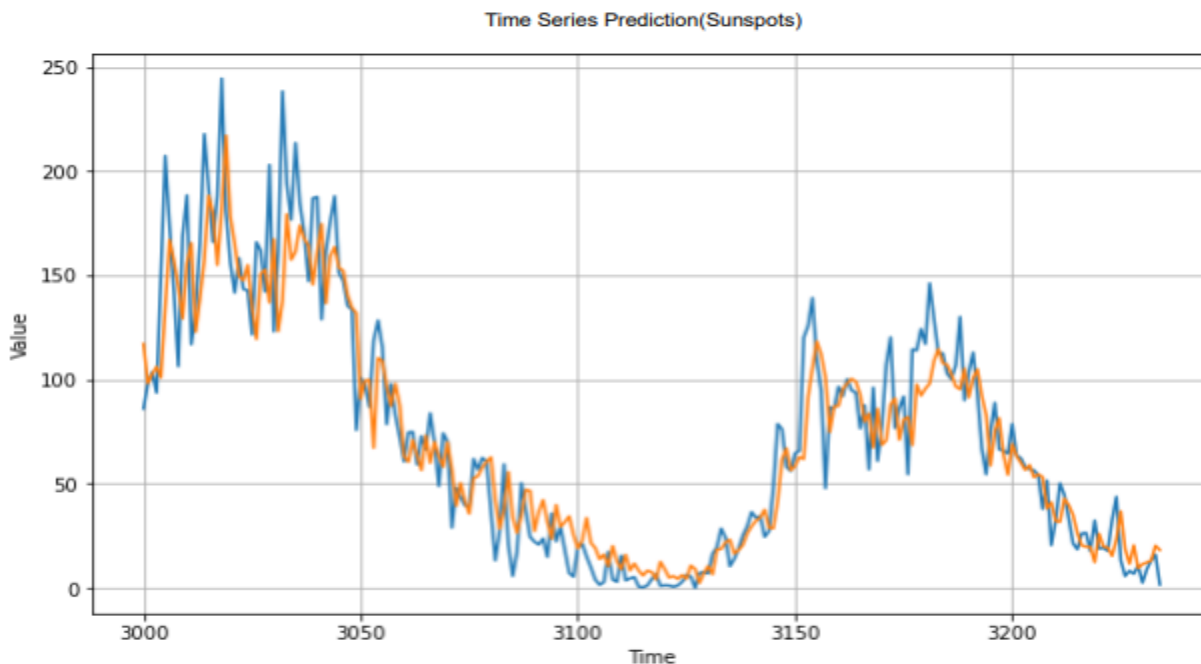
43.33463829787234

Forecasting Using Linear Regression [10]

We can create a Linear Regression model by using a single neuron Deep Neural Network.

The input shape we fed into the neural network will be equal to window size, the output will be a single value.

We can also see the resultant weights and the bias value by using `get_weights()` method.



Mean Absoule Error for Linear Regression(single neuron)


```
tf.keras.metrics.mean_absolute_error(x_valid, results).numpy()
```

15.3134165

This is the least mean absolute error until now, so from this we can conclude that Linear Regression seems to perform better the Naïve & Moving Averages approach.

Forecasting Using DNN (Deep Neural Networks) [6]

Creating a DNN model is similar to creating a Linear Regression model, the difference is that we will be using multiple layers also called DNN layers in our model, for this project we will be implementing 3 layer mode, 2 DNN layers and 1 output layer, We will be using **relu** activation function.



Mean Square Error.

```
tf.keras.metrics.mean_absolute_error(x_valid, results).numpy()
```

15.176499

Forecasting Using RNNs & Convolutions [6] [8]


A recurrent neural network (RNN) is a type of artificial neural network which uses sequential data or time series data.

Like feedforward and convolutional neural networks (CNNs), recurrent neural networks utilize training data to learn. They are distinguished by their “memory” as they take information from prior inputs to influence the current input and output. [9]

In our model we have used 2 RNN layers along with convolution layer to forecast our time series.

We'll first run for 100 epochs and determine optimal learning rate for our stochastic gradient descent (SGD), and then we retrain our model for 500 epochs to get our **MAE (Mean Absolute Error)** and predictions.

RNNs have a memory, but the effect of earlier inputs diminish over time in RNNs unlike LSTMs.



Mean Square Error.

```
# MAE  
tf.keras.metrics.mean_absolute_error(x_valid, rnn_forecast).numpy()
```

15.583032

Forecasting Using LSTM & Convolutions [6][8][11]


For our final model we will be using convolutions along with LSTMS, this helps us to create a model that fits really well to our data, we can also use various tuning methods to further improve our performance.

LSTMS also known as Long Short Term Memory, It is a type of RNN(Recurrent Neural Network) which helps us preserve the information about current state throughout the network, and thus this helps us make better predictions as recent past values have more impact on future values.

In this model, we used one convolution layer with 32 filters, after that we have added 2 LSTM layers, after that we have added 3 dense neuron layers, with 30, 10 and 1 units respectively.


First we train the model for 100 epochs to get an optimal learning rate for our stochastic gradient descent (SGD) optimizer.

(1e-08, 0.0001, 0.0, 60.0)



After getting the optimal learning rate, we retrain the model for 500 epochs with our optimal learning rate.

After creating the model, we will create a forecast array of predicted values, and then we plot our results.

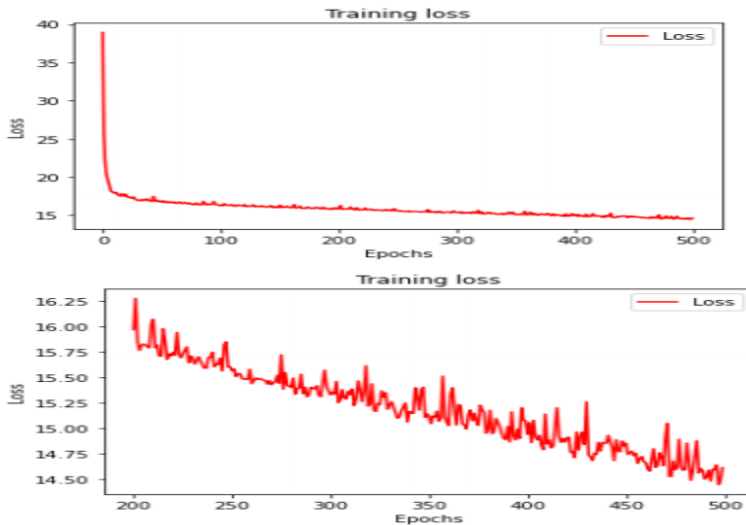


Mean Square Error.

```
tf.keras.metrics.mean_absolute_error(x_valid, rnn_forecast).numpy()
```

15.089574

Training Loss vs Epochs.



Forecasted Values.

```
print(rnn_forecast)
```

```
[135.92502    99.271904   105.50781    129.8834    109.807526   150.31667
 192.19649    168.62233    179.26598    126.94186    148.1779    169.16774
 158.67053    159.65434    195.03677    190.81548    185.02423    192.65051
 177.28357    201.87624    190.02647    177.29008    172.10686    182.73575
 159.40915    149.92252    153.19821    164.56311    159.92856    134.9303
 199.1364     128.0686     148.98816    218.37451    172.10358    183.2
 178.08499    152.958      152.39517    164.84302    174.30031    181.5857
 155.59616    163.76837    187.73816    164.62791    160.91759    153.94258
 147.61807    142.00415    104.43032    113.7124    108.309616    95.00554
 112.304596   115.40711    100.481155    90.23153     88.875725    77.811775
.....
```

As we see that **Mean absolute Error** is quite low, thus indicating that with little more tuning, we can predict sunspots activity quite accurately.

LSTMs performance turns out to be better the RNNs, contributing to the fact that LSTMs retains the effect of earlier inputs throughout the state while RNNs do not.

CODE (LSTM AND CONVOLUTIONS)

```
import tensorflow as tf
print(tf.__version__)

import numpy as np
import matplotlib.pyplot as plt
def plot_series(time, series, format="-", start=0, end=None):
    plt.plot(time[start:end], series[start:end], format)
    plt.xlabel("Time")
    plt.ylabel("Value")
    plt.grid(True)
```


2.3.0

Importing Data CSV file and visualizing

```
import csv
time_step = []
sunspots = []

with open('Sunspots.csv') as csvfile:
    reader = csv.reader(csvfile, delimiter=',')
    next(reader)
    for row in reader:
        sunspots.append(float(row[2]))
        time_step.append(int(row[0]))
```

```
series = np.array(sunspots)
time = np.array(time_step)
plt.figure(figsize=(10, 6))
plot_series(time, series)
```



#Creating new window size, batch size and shuffle buffer size

```
split_time = 3000
time_train = time[:split_time]
x_train = series[:split_time]
time_valid = time[split_time:]
x_valid = series[split_time:]
```

```
window_size = 30
batch_size = 32
shuffle_buffer_size = 1000
```

#Converting the dataset into windowed dataset

```
def windowed_dataset(series, window_size, batch_size, shuffle_buffer):
    series = tf.expand_dims(series, axis=-1)
    ds = tf.data.Dataset.from_tensor_slices(series)
    ds = ds.window(window_size + 1, shift=1, drop_remainder=True)
    ds = ds.flat_map(lambda w: w.batch(window_size + 1))
    ds = ds.shuffle(shuffle_buffer)
    ds = ds.map(lambda w: (w[:-1], w[1:]))
    return ds.batch(batch_size).prefetch(1)
```

#Helper Function for forecasting the data.

```
def model_forecast(model, series, window_size):
    ds = tf.data.Dataset.from_tensor_slices(series)
    ds = ds.window(window_size, shift=1, drop_remainder=True)
    ds = ds.flat_map(lambda w: w.batch(window_size))
    ds = ds.batch(32).prefetch(1)
    forecast = model.predict(ds)
    return forecast
```

```
tf.keras.backend.clear_session()
window_size = 64
batch_size = 256
train_set = windowed_dataset(x_train, window_size, batch_size, shuffle_buffer_size)
print(train_set)
print(x_train.shape)
```

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Conv1D(filters=32, kernel_size=5,
                           strides=1, padding="causal",
                           activation="relu",
                           input_shape=[None, 1]),
    tf.keras.layers.LSTM(64, return_sequences=True),
    tf.keras.layers.LSTM(64, return_sequences=True),
    tf.keras.layers.Dense(30, activation="relu"),
    tf.keras.layers.Dense(10, activation="relu"),
    tf.keras.layers.Dense(1),
    tf.keras.layers.Lambda(lambda x: x * 400)
])
```

```
lr_schedule = tf.keras.callbacks.LearningRateScheduler(
    lambda epoch: 1e-8 * 10**(epoch / 20))
optimizer = tf.keras.optimizers.SGD(lr=1e-8, momentum=0.9)
model.compile(loss=tf.keras.losses.Huber(),
              optimizer=optimizer,
              metrics=["mae"])
history = model.fit(train_set, epochs=100, callbacks=[lr_schedule])
```

```

WARNING:tensorflow:AutoGraph could not transform <function windowed_dataset.<locals>.<lambda> at 0x0000022F7C483040> and will run it as-is.
Please report this to the TensorFlow team. When filing the bug, set the verbosity to 10 (on Linux, `export AUTOGRAPH_VERBOSITY=10`) and attach the full output.
Cause: module 'gast' has no attribute 'Index'
To silence this warning, decorate the function with @tf.autograph.experimental.do_not_convert
WARNING: AutoGraph could not transform <function windowed_dataset.<locals>.<lambda> at 0x0000022F7C483040> and will run it as-is.
Please report this to the TensorFlow team. When filing the bug, set the verbosity to 10 (on Linux, `export AUTOGRAPH_VERBOSITY=10`) and attach the full output.
Cause: module 'gast' has no attribute 'Index'
To silence this warning, decorate the function with @tf.autograph.experimental.do_not_convert
<PrefetchDataset shapes: ((None, None, 1), (None, None, 1)), types: (tf.float64, tf.float64)>
(3000,)
Epoch 1/100
12/12 [=====] - 2s 199ms/step - loss: 79.8340 - mae: 80.3314
Epoch 2/100
12/12 [=====] - 2s 181ms/step - loss: 78.0944 - mae: 78.5918
Epoch 3/100
12/12 [=====] - 2s 185ms/step - loss: 75.4519 - mae: 75.9497
Epoch 4/100
12/12 [=====] - 2s 188ms/step - loss: 72.2678 - mae: 72.7658
Epoch 5/100
12/12 [=====] - 2s 183ms/step - loss: 68.7693 - mae: 69.2672
Epoch 6/100


```

```

plt.semilogx(history.history["lr"], history.history["loss"])
plt.axis([1e-8, 1e-4, 0, 60])

```

```
(1e-08, 0.0001, 0.0, 60.0)
```



```

tf.keras.backend.clear_session()
train_set = windowed_dataset(x_train, window_size=60, batch_size=100, shuffle_buffer=shuffle_buffer_size)
model = tf.keras.models.Sequential([
    tf.keras.layers.Conv1D(filters=60, kernel_size=5,
                           strides=1, padding="causal",
                           activation="relu",
                           input_shape=[None, 1]),
    tf.keras.layers.LSTM(60, return_sequences=True),
    tf.keras.layers.LSTM(60, return_sequences=True),
    tf.keras.layers.Dense(30, activation="relu"),
    tf.keras.layers.Dense(10, activation="relu"),
    tf.keras.layers.Dense(1),
    tf.keras.layers.Lambda(lambda x: x * 400)
])

optimizer = tf.keras.optimizers.SGD(lr=1e-5, momentum=0.9)
model.compile(loss=tf.keras.losses.Huber(),
              optimizer=optimizer,
              metrics=["mae"])
history = model.fit(train_set, epochs=500)

```

WARNING:tensorflow:AutoGraph could not transform <function windowed_dataset.<locals>.<lambda> at 0x0000022F0E910280> and will run it as-is.
Please report this to the TensorFlow team. When filing the bug, set the verbosity to 10 (on Linux, `export AUTOGRAPH_VERBOSITY=10`) and attach the full output.

Cause: module 'gast' has no attribute 'Index'

To silence this warning, decorate the function with @tf.autograph.experimental.do_not_convert

WARNING: AutoGraph could not transform <function windowed_dataset.<locals>.<lambda> at 0x0000022F0E910280> and will run it as-is.

Please report this to the TensorFlow team. When filing the bug, set the verbosity to 10 (on Linux, `export AUTOGRAPH_VERBOSITY=10`) and attach the full output.

Cause: module 'gast' has no attribute 'Index'

To silence this warning, decorate the function with @tf.autograph.experimental.do_not_convert

Epoch 1/500

30/30 [=====] - 4s 137ms/step - loss: 38.9182 - mae: 39.4146

Epoch 2/500

30/30 [=====] - 4s 119ms/step - loss: 25.7634 - mae: 26.2578

Epoch 3/500

30/30 [=====] - 4s 123ms/step - loss: 22.0849 - mae: 22.5786

Epoch 4/500

30/30 [=====] - 3s 112ms/step - loss: 20.4691 - mae: 20.9621

Epoch 5/500

30/30 [=====] - 4s 123ms/step - loss: 19.7838 - mae: 20.2763

Epoch 6/500

30/30 [=====] - 4s 121ms/step - loss: 19.2844 - mae: 19.7767

Epoch 7/500

```


rnn_forecast = model_forecast(model, series[...], np.newaxis], window_size)
rnn_forecast = rnn_forecast[split_time - window_size:-1, -1, 0]

```

```

plt.figure(figsize=(10, 6))
plot_series(time_valid, x_valid)
plot_series(time_valid, rnn_forecast)

```



```
tf.keras.metrics.mean_absolute_error(x_valid, rnn_forecast).numpy()
```

15.089574

```
import matplotlib.image as mpimg
import matplotlib.pyplot as plt

#-----
# Retrieve a list of list results on training and test data
# sets for each training epoch
#-----
loss=history.history['loss']

epochs=range(len(loss)) # Get number of epochs

#-----
# Plot training and validation loss per epoch
#-----
plt.plot(epochs, loss, 'r')
plt.title('Training loss')
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend(["Loss"])


plt.figure()
```

```
zoomed_loss = loss[200:]
zoomed_epochs = range(200,500)

#-----
# Plot training and validation loss per epoch
#-----
plt.plot(zoomed_epochs, zoomed_loss, 'r')
plt.title('Training loss')
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend(["Loss"])

plt.figure()
```

<Figure size 432x288 with 0 Axes>



```
print(rnn_forecast)
```

```
[135.92502  99.271904 105.50781 129.8834 109.807526 150.31667
192.19649 168.62233 179.26598 126.94186 148.1779 169.16774
158.67053 159.65434 195.03677 190.81548 185.02423 192.65051
177.28357 201.87624 190.02647 177.29008 172.10686 182.73575
159.40915 149.92252 153.19821 164.56311 159.92856 134.9303
199.1364 128.0686 148.98816 218.37451 172.10358 183.2
178.08499 152.958 152.39517 164.84302 174.30031 181.5857
155.59616 163.76837 187.73816 164.62791 160.91759 153.94258
147.61807 142.00415 104.43032 113.7124 108.309616 95.00554
112.304596 115.40711 100.481155 90.23153 88.875725 77.811775
71.621796 73.536644 77.760765 85.392105 72.36296 69.2541
69.323746 74.63767 76.34568 55.404366 64.97453 76.79944
50.71929 50.283485 46.622158 34.97485 36.70865 48.570408
47.02358 52.555763 60.58049 34.326183 29.493862 27.580122
43.976254 20.87406 17.527748 16.793663 14.4772415 14.254313
14.442165 17.74445 18.561777 22.075977 21.341286 24.885897
22.398823 23.092352 19.850086 13.912702 12.291941 11.018311
13.696872 12.136144 12.551162 11.527196 10.368964 9.625668
```

.....

BIBLIOGRAPHY/ REFERENCES

This project could not be possible without the help of online resources. Some of the resources includes:-

1. GitHub Resources and third party libraries.
2. Online videos & Tutorials.
3. Online QnA forums (Stackoverflow, quora etc.)
4. Babcock H.W., "The topology of the Sun's magnetic field and the 22-Year cycle", *Astrophysical Journal*, 2, pp. 572–587, 1961.
5. R. P. Kane, "An estimate for the size of sunspot cycle 24," *Solar Physics*, vol. 282, no. 1, pp. 87–90, 2013. View at: [Publisher Site](#) | [Google Scholar](#)
6. https://www.tensorflow.org/tutorials/structured_data/time_series
7. <https://www.analyticsvidhya.com/blog/2018/02/time-series-forecasting-methods/>
8. <https://towardsdatascience.com/prediction-and-analysis-of-time-series-data-using-tensorflow-2136ef633018>
9. <https://www.ibm.com/cloud/learn/recurrent-neural-networks>
10. <https://www.analyticsvidhya.com/blog/2021/06/linear-regression-using-neural-networks/>
11. <https://machinelearningmastery.com/gentle-introduction-long-short-term-memory-networks-experts/>

FUTUREWORK

We will try to add some extra functionality and apply some parameter tuning to the model and improve the accuracy.

We will also be testing with bidirectional LSTMS and see how much our accuracy improves.

Our goal is to get the **mean absolute error** value below 10.00.

THANK YOU